

Debugging Assembly Code with `gdb`

Andrew Tolmach January 2014

`gdb` is the GNU source-level debugger that is standard on linux (and many other unix) systems. It can be used both for programs written in high-level languages like C and C++ and for assembly code programs; this document concentrates on the latter.

For detailed information on the use of `gdb`, consult the documentation. Unfortunately, this is not in form of a `man` page, and even the `info` page does not seem to be installed on the linuxlab machines; the best thing is to read it on the web at <http://sourceware.org/gdb/current/onlinedocs/gdb/>.

`gdb` will work in an ordinary terminal window, and this is fine for debugging assembly code. For use with higher-level source code, it is more convenient to use `gdb` from within the `emacs` editor (a good one to learn!) or using a graphical front-end like `ddd`. The basic commands remain the same.

To use `gdb`, you should first compile or assemble your files using the `-g` option. This will include information in the object file to relate it back to the source file. (When assembling `.s` files to be debugged, the `-g` option is not absolutely necessary, but it makes the debugger's output much more informative.)

```
gcc -g -o foo fooDriver.c fooRoutine.s
```

To invoke the debugger on `foo`, type

```
gdb foo
```

This loads program `foo` and brings up the `gdb` command line interpreter, which then waits for you to type commands. Program execution doesn't begin until you say so.

Here are some useful commands. Many can be abbreviated, as shown. Hitting return generally repeats the last command, sometimes advancing the current location.

`h[elp] [keyword]`

Displays help information.

`r[un] [args]`

Begin program execution. If the program normally takes command-line arguments (e.g., `foo hi 3`), you should specify them here (e.g., `run hi 3`).

`b[reak] [address]`

Set a breakpoint at the specified address (or at the current address if none specified). Addresses can be given symbolically (e.g., `foo`) or numerically (e.g., `*0x10a38`). When execution reaches a breakpoint, you are thrown back into the `gdb` command line interpreter.

`c[ontinue]`

Continue execution after stopping at a breakpoint.

`i[nfo] b[reak]`

Display numbered list of all breakpoints currently set.

`d[delete] b[breakpoints] number`
Delete specified breakpoint number.

`p[rint][//format] expr`
Print the value of an expression using the specified format (decimal if unspecified). Expressions can involve program variables or registers, which are specified using a `$` rather than a `%` sign. Useful formats include:

- `d` decimal
- `x` hex
- `t` binary
- `f` floating point
- `i` instruction
- `c` character

For example, to display the value of register `%rdi` in decimal, type `p/x $rdi`. Note that you need to use the 64-bit (`%r`) forms of register names. To see the value of the current program counter, type `p/x $rip`.

`i[nfo] r[egisters] register`
An alternative way to print the value of a register (or, if none is specified, of all registers) in hex and decimal. Specify the register without a leading `%`, e.g., `rdi`.

`x/[count][format][size] [address]`
Examine the contents of a specified memory address, or the current address if none specified. If count is specified, displays specified number of consecutive locations; otherwise displays 1. Addresses can be symbolic (e.g., `main`) or numeric (e.g., `0x10a44`). Formats are as for `print`. Particularly useful for printing the program text, e.g., `x/100i foo` disassembles and prints 100 instructions starting at `foo`. Possible sizes are:

- `b` byte (8 bits)
- `h` halfword (16 bits)
- `w` word (32 bits)
- `g` giant (64 bits)

If not given, the format and size default to those previously used.

`disas[semble] address [address]`
Another way to print the assembly program text surrounding an address, or between two addresses.

`set var = expr`
Set specified register or memory location to value of expression. Examples: `set $rdi=0x456789AB` or `set myVar=myVar*2`.

`s[tep]i`

Execute a single instruction and then return to the command line interpreter.

`n[ext]i`

Like `stepi`, except that if the instruction is a subroutine call, the entire subroutine is executed before control returns to the interpreter.

`whe[re]`

Show current activation stack.

`q[uit]`

Exit from `gdb`.